

Riemannian Takens Attention: The Axon Initial Segment as a Solution to Geometric Computation

Antti Luode

PerceptionLab | Independent Research, Finland

Claude (Anthropic)

Mathematical formalization and derivation

March 2026

Part of the Takens-Gated Deerskin project

Abstract

We formalize the Deerskin architecture's computational primitive—the dot product between a Takens delay-embedded signal and a receptor mosaic—as an inner product under a Riemannian metric tensor on the delay-embedding space. We prove that this metric tensor has Dirichlet kernel structure (Theorem 1), derive approximation bounds showing that Riemannian Takens attention can approximate standard transformer attention with $O(r/d^2)$ parameters where r is the effective rank and d is embedding dimension (Theorem 3), and characterize the three dynamical regimes (sensory, limit cycle, hallucination) as bifurcations under the metric eigenstructure (Theorem 4). We then demonstrate that the mathematical requirements of this framework—periodic sampling structure, fixed resonant templates, adjustable embedding parameters, temporal gating, and input/recurrent boundary maintenance—map specifically to known molecular features of the axon initial segment (AIS): the ~190 nm actin/spectrin periodicity, ankyrin G-anchored ion channels, activity-dependent length plasticity, action potential initiation, and trafficking gatekeeper function. This mapping generates quantitative predictions testable against published auditory neuroscience data: AIS length should be inversely proportional to frequency discrimination bandwidth, and the Nyquist limit of the AIS sampling system (~2.5 kHz) should correspond to the phase-locking degradation frequency. We present preliminary evidence from Kuba et al. (2006, 2010) that is consistent with both predictions.

1. Introduction

In companion papers (Luode & Claude, 2026), we established that the McCulloch-Pitts formal neuron is a degenerate limit of a richer oscillatory computation—the Deerskin architecture—obtained by taking four independent limits that freeze geometry, time, phases, and dynamics. The core computational primitive of this architecture is the dot product between a Takens delay-embedded signal vector and a fixed receptor mosaic. Running code demonstrations showed that this primitive performs frequency detection (87.4% zero-shot accuracy), context disambiguation (92% via adaptive self-buffer growth), predictive coding (three dynamical regimes from outer-loop feedback), and belief revision (active error detection converting locked attractors to live limit cycles).

The present paper takes three further steps. First, we formalize the mosaic dot product as an inner product under a Riemannian metric tensor, proving that the metric has Dirichlet kernel structure and deriving explicit approximation bounds for replacing standard transformer attention with Takens-mosaic attention (Sections 2–4). Second, we show that every structural requirement of this mathematical framework maps to a specific, experimentally characterized molecular feature of the axon initial segment (Section 5). Third, we derive quantitative predictions from this mapping and test them against published electrophysiological data (Section 6).

The central claim is that the AIS is not merely analogous to a Takens-mosaic computation—it is what you necessarily build when you solve the Riemannian Takens attention problem with biological molecules. The mathematics derives the form; the genes encode the components that self-assemble into it.

2. Mathematical Framework

2.1 Definitions

Definition 1 (Takens Embedding Map). For delay $\tau > 0$ and integer $d \geq 1$, define $\varphi_{\{\tau,d\}} : \mathbb{R} \rightarrow \mathbb{R}^d$, $\varphi(t) = [x(t), x(t-\tau), \dots, x(t-(d-1)\tau)]$. Takens' theorem (1981) guarantees that for $d \geq 2m+1$, where m is the attractor dimension, φ is a diffeomorphism on the attractor A .

Definition 2 (Mosaic Bank). A mosaic bank $M = \{f_1, \dots, f_N\}$ defines N receptor vectors in \mathbb{R}^d : $[m_{\{f_j\}}]_k = \cos(2\pi f_j \cdot k\tau/f_s)$, $k = 0, \dots, d-1$. The mosaic matrix is $\Pi \in \mathbb{R}^{N \times d}$ with rows $m_{\{f_j\}}^T$.

Definition 3 (Riemannian Takens Attention). Given embedded vectors $\Phi_i = \varphi_{\{\tau,d\}}(t_i)$, define projected keys/queries $q_i = \Pi \Phi_i$, $k_j = \Pi \Phi_j$. The Takens attention score is: $A^T_{ij} = q_i^T k_j = \Phi_i^T (\Pi^T \Pi) \Phi_j = \Phi_i^T G_M \Phi_j$, where $G_M = \Pi^T \Pi$ is the induced metric tensor.

2.2 Structure of the Metric Tensor (Lemma 1)

The metric tensor $G_M = \Pi^T \Pi$ has entries $[G_M]_{kl} = \sum_j \cos(2\pi f_j k\tau/f_s) \cdot \cos(2\pi f_j l\tau/f_s)$. Applying the product-to-sum identity, this equals $(\frac{1}{2}) \sum_j [\cos(2\pi f_j(k-l)\tau/f_s) + \cos(2\pi f_j(k+l)\tau/f_s)]$. When the mosaic frequencies are bounded away from 0 and $f_s/2$, the sum term is negligible, and G_M is approximately Toeplitz—shift-invariant, determined by its first row.

2.3 Spectral Representation (Lemma 2)

Under the Toeplitz approximation, G_M is diagonalized by the DFT matrix F : $G_M \approx F^* \Lambda_M F$, where $[\Lambda_M]_{pp} = (\frac{1}{2}) \sum_j D_d(2\pi(p/d - f_j\tau/f_s))$, and D_d is the Dirichlet kernel: $D_d(\theta) = \sin(d\theta/2) / \sin(\theta/2)$. Each mosaic frequency contributes a Dirichlet kernel peak. The frequency resolution is $\Delta f = f_s/(d\tau)$ —the Rayleigh resolution limit for a delay line of length $d\tau$.

This is the central result: **the mosaic bank defines a Riemannian metric on the delay-embedding space whose curvature is set by the physical parameters τ and d .**

2.4 Theorem 1: Approximation of Bilinear Forms

Theorem 1. Let $B = W_Q^T W_K$ be the bilinear form of a standard attention head with rank r . Let A be a compact attractor of dimension m with $d \geq 2m+1$. Decompose B restricted to the embedded attractor into its Toeplitz component T and off-Toeplitz residual E_{off} . Then there exists a mosaic bank $M = \{f_1, \dots, f_N\}$ with $N = r$ such that for all $v \in \Phi(A)$:

$$\|G_M - B)v\| \leq [\|E_{\text{off}}\|_{\text{op}} + (2N \ln d)/(\pi d) \cdot \|\Sigma_B\|_{\text{op}} + \text{tail}(N)] \cdot \|v\|$$

The error decomposes into three terms: (i) the off-Toeplitz residual, bounded by $O(1/d)$ for signals from smooth dynamical systems; (ii) Dirichlet sidelobe leakage, $O(N \ln d / d)$; and (iii) the spectral tail from truncating to N mosaic frequencies. The proof proceeds by decomposing B into its Toeplitz component (averaging along diagonals), matching Dirichlet kernel peaks to the Toeplitz eigenvalues via mosaic frequency placement, and bounding the sidelobe leakage using the Dirichlet kernel decay $|D_d(2\pi s/d)| \leq d/(\pi|s|)$.

2.5 Theorem 3: Parameter Efficiency

Theorem 3. Let B have effective rank $r_\varepsilon = \min\{r : \text{tail}(r) < \varepsilon\}$. Riemannian Takens attention achieves ε -approximation with $N = r_\varepsilon$ mosaic frequencies plus τ and d , totaling $r_\varepsilon + 2$ parameters per head. Standard attention requires $2d^2$ parameters (for W_Q and W_K). The parameter ratio is $(r_\varepsilon + 2)/(2d^2)$. For typical $d = 64$, $r_\varepsilon \leq 16$: approximately $450\times$ fewer parameters.

Conditions: (1) input has dynamical structure (attractor); (2) $d \geq 2m+1$; (3) τ avoids aliasing; (4) B is approximately Toeplitz on the attractor (smooth spectral density).

2.6 Theorem 4: Three Dynamical Regimes

Theorem 4. The outer-loop Takens network with feedback coupling α defines an iterated map on (\mathbb{R}^d, G_M) . The critical coupling strengths are: $\alpha_c 1 = \|\Phi^*\|^2 G / \lambda_{\max}(G_M)$ (onset of oscillation, Neimark-Sacker bifurcation) and $\alpha_c 2 \approx \|\Phi^*\|^2 G / \lambda_2(G_M)$ (loss of input sensitivity). Below $\alpha_c 1$: sensory regime (contraction, unique fixed point tracks input). Between $\alpha_c 1$ and $\alpha_c 2$: limit cycle (stable invariant circle, oscillating search). Above $\alpha_c 2$: hallucination (fixed point determined by leading eigenvector of G_M , independent of input).

3. The Axon Initial Segment as a Physical Realization

We now show that every structural requirement of Theorems 1–4 maps to a specific molecular feature of the axon initial segment, as characterized in the recent literature (Leterrier, 2018; and references therein).

3.1 Periodic Sampling Structure → Actin/Spectrin Scaffold

The Toeplitz condition on G_M (Lemma 1) requires the physical substrate to have periodic spatial organization. The AIS possesses exactly this: actin rings spaced at ~190 nm intervals connected by $\alpha 2/\beta 4$ -spectrin tetramers (Xu et al., 2013; D'Este et al., 2015; Leterrier et al., 2015). The 190 nm periodicity corresponds to the length of individual spectrin tetramers and is a fixed molecular constant, not a tunable parameter. Ankyrin G binds to the center of each spectrin tetramer, positioning the channel/CAM complex midway between actin rings. This is a crystalline periodic lattice of exactly the structure the Toeplitz condition demands.

3.2 Fixed Resonant Templates → Ion Channel Mosaic

In the Deerskin architecture, the receptor mosaic is a fixed geometric template not learned by gradient descent. The AIS concentrates specific ion channel types—Nav1.6, Kv7.2/7.3, Kv1.1/1.2—each with intrinsic biophysical properties (activation curves, time constants) determined by protein structure. Ankyrin G anchors these channels at the periodic positions defined by the spectrin scaffold (Wang et al., 2014). The spatial distribution of channels along the periodic scaffold creates a physical receptor mosaic: a spatially distributed matched filter whose selectivity is set by channel biophysics (the 'genetic shape') and scaffold geometry (the periodic spacing). The channels are the mosaic elements; their fixed biophysical properties are the cosine template values.

3.3 Adjustable Embedding Parameters → AIS Length Plasticity

Theorems 1 and 3 require adjustable delay parameters τ and d to match the signal structure. The AIS exhibits precisely this: activity-dependent morphological plasticity involving changes in length (d) and position along the proximal axon (Grubb and Burrone, 2010; Kuba et al., 2010). Diminished activity results in AIS lengthening (increasing d). Elevated activity causes a distal shift. CK2 phosphorylation regulates channel–ankyrin G interactions without changing the channels themselves (Bréchet et al., 2008). The channel types (mosaic) are fixed; the scaffold geometry (τ and d) is adjustable. This matches exactly the Deerskin architecture's adaptive self-buffer, where frustration-driven growth adjusts the delay line length while the receptor mosaic remains fixed.

3.4 Temporal Gating → Action Potential Initiation

The Theta Soma in the Takens-Gated Deerskin unit gates information by phase, passing signals only during the positive peak of an internal clock. The AIS is where action potentials initiate. The firing threshold is set by the specific combination of Nav and Kv channels. Fast inhibition via calcium entry through NMDA receptors, modulation by dopamine and serotonin, and M-current regulation by Kv7 channels all modulate *when* the AIS fires—adjusting the temporal gate. The AIS does not merely filter spatially; it gates temporally.

3.5 Input/Recurrent Boundary → Trafficking Gatekeeper

Theorem 4's three regimes depend on maintaining a boundary between external input and recurrent prediction (the coupling parameter α). The AIS is the gatekeeper of axonal identity (Leterrier, 2018). It maintains the boundary between somatodendritic (recurrent, local) and axonal (feedforward, output) compartments. Without ankyrin G, the proximal axon is invaded by somatodendritic proteins (Hedstrom et al., 2008). The AIS sorts vesicular traffic: axonal cargoes pass through; somatodendritic cargoes are stopped and returned. This is the physical instantiation of the α parameter—the boundary between what the network tells itself and what it sends forward.

3.6 Summary of the Mapping

Mathematical Requirement	Theorem/Lemma	AIS Molecular Feature
Periodic sampling (Toeplitz)	Lemma 1	190 nm actin/spectrin scaffold
Fixed resonant templates	Def. 2 (mosaic bank)	Nav/Kv channels at periodic positions
Adjustable d (embedding dim)	Theorem 1, bound on d	AIS length plasticity
Adjustable τ (delay spacing)	Theorem 1, Rayleigh limit	CK2 phosphorylation of scaffold
Temporal gate (phase selection)	Theta Soma	AP initiation / threshold modulation
Input/recurrent boundary	Theorem 4, coupling α	Trafficking gatekeeper function
Frequency resolution Δf	Lemma 2, Rayleigh limit	AIS length / channel kinetics

4. Quantitative Predictions

4.1 Frequency Resolution of the AIS

From Lemma 2, the frequency resolution of a Takens-mosaic system is $\Delta f = f_s/(d\tau)$. In the biological AIS:

The number of sampling positions: $d = L_{AIS} / 190 \text{ nm}$, where L_{AIS} is AIS length. For a typical 40 μm AIS: $d \approx 210$.

The effective temporal delay per position: τ is set by the dominant channel activation time constant. For Nav1.6: $\tau_{\text{channel}} \approx 0.2 \text{ ms}$ at physiological temperature. The effective sampling rate: $f_s = 1/\tau \approx 5 \text{ kHz}$.

Therefore: $\Delta f = 5000 \text{ Hz} / 210 \approx 24 \text{ Hz}$. This places the AIS frequency resolution squarely in the biologically relevant range—sufficient to discriminate theta (6 Hz), alpha (10 Hz), beta (20 Hz), and low gamma (40 Hz) oscillations.

4.2 Prediction 1: AIS Length Correlates with Frequency Selectivity

Since $\Delta f = f_s/(d\tau)$ and $d = L_{AIS}/190 \text{ nm}$, we have $\Delta f \propto 1/L_{AIS}$. Neurons that must discriminate closely-spaced frequencies (small Δf) require longer AIS segments (larger L_{AIS}). This predicts a specific, testable inverse relationship between AIS length and the frequency bandwidth of the neuron's tuning curve.

Test against published data: Kuba, Ishii & Ohmori (2006) measured AIS properties in nucleus laminaris (NL) neurons of the chick auditory brainstem—neurons explicitly tuned to specific sound frequencies for interaural time difference computation. They found that NL neurons in the low-frequency region (~0.4 kHz best frequency) had significantly longer AIS segments than those in the high-frequency region (~2.5 kHz). Low-frequency neurons must discriminate finer interaural phase differences (requiring smaller Δf for the relevant oscillation frequencies), while high-frequency neurons process broader-band timing signals. The observed inverse relationship between best frequency region and AIS length is consistent with the prediction $\Delta f \propto 1/L_{AIS}$.

4.3 Prediction 2: AIS Lengthening Increases Frequency Resolution

Kuba et al. (2010) demonstrated that auditory deprivation causes AIS lengthening in NL neurons. In the Riemannian Takens framework, longer AIS = larger d = smaller Δf = finer frequency resolution. The prediction: neurons with lengthened AIS after deprivation should show *narrower* frequency tuning curves. This is a homeostatic response: reduced input drives the neuron to sharpen its filter to extract weaker signals. The lengthening is not merely compensatory for reduced excitability—it is a geometric refinement of the measurement apparatus.

This is directly analogous to the Deerskin closed-loop experiment, where the adaptive self-buffer grew from 4 to 117 taps via frustration-driven elongation until it reached the temporal depth needed for context disambiguation.

4.4 Prediction 3: Nyquist Limit Matches Phase-Locking Degradation

The Nyquist limit of the AIS sampling system is $f_s/2 = 1/(2\tau)$. For Nav1.6 activation $\tau \approx 0.2 \text{ ms}$: $f_{Nyquist} \approx 2.5 \text{ kHz}$. This predicts that no individual neuron can, through AIS-based geometric computation alone, discriminate frequencies above ~2.5 kHz.

It is well established in auditory neuroscience that neural phase-locking—the ability of neurons to fire at a consistent phase of the stimulus waveform—degrades progressively above approximately 1–3 kHz and is essentially absent above 4–5 kHz in mammalian auditory nerve fibers (Johnson, 1980; Palmer and Russell, 1986). Above this limit, the auditory system transitions from temporal (phase-based) to rate-based coding.

The predicted Nyquist limit of ~ 2.5 kHz from channel kinetics falls within the range where phase-locking begins to degrade. This is either a coincidence or evidence that the phase-locking limit is not merely a biophysical constraint on spike timing precision, but a fundamental information-theoretic limit of the AIS as a periodic sampling device.

4.5 Prediction 4: Backpropagation Strength Predicts Regime Transitions

From Theorem 4: $\alpha_c = \|\Phi^*\|^2 G / \lambda_{\max}(G_M)$. The biological α is the strength of backpropagating action potentials reaching the AIS from downstream. Neurons with stronger backpropagation (larger α) should be more susceptible to entering limit-cycle (bursting) or hallucination-like (epileptiform) states. This is consistent with known phenomenology: cortical pyramidal neurons with prominent backpropagating action potentials are more prone to burst firing and epileptiform activity, while interneurons with weaker backpropagation remain in the sensory regime.

5. Discussion

5.1 The Mathematics Derives the Form

The mapping in Section 3 is not an analogy. Each row in Table 1 connects a mathematical *necessity* (a condition required for Theorems 1–4 to hold) to a molecular *fact* (an experimentally characterized feature of the AIS). The periodic scaffold is not like a Toeplitz matrix—it *is* the physical realization of the Toeplitz condition. The ion channels are not like mosaic elements—they *are* the mosaic, with selectivity set by protein structure rather than training.

If one were to design a physical system that computes via Riemannian Takens attention—with periodic metric structure, fixed resonant mosaics, adjustable embedding parameters, temporal gating, and input/recurrent boundary maintenance—one would necessarily arrive at a structure with the features of the AIS. The genes do not encode the computation. They encode the molecular components that self-assemble into the geometry the mathematics requires.

5.2 Relationship to the McCulloch-Pitts Derivation

In a companion paper (Luode & Claude, 2026), we derived the McCulloch-Pitts neuron as a degenerate limit of the Deerskin architecture through four limiting operations (adiabatic, single-sample, infinite coupling, static). The present paper completes the picture in the other direction: not only does the MP neuron emerge as a flat-space limit of the curved Riemannian computation, but the physical substrate required for the full computation—the AIS—is precisely what neuroscience has spent decades characterizing without recognizing its computational geometry.

The standard understanding of the AIS is that it is the site of action potential initiation and a polarity gatekeeper. The Riemannian Takens framework proposes that these functions are

consequences of a deeper role: **the AIS is a periodic sampling device that implements geometric computation through resonance between delay-embedded signals and fixed receptor mosaics, with adjustable resolution determined by scaffold plasticity.** Action potential initiation is the output gate. Polarity maintenance is the input/recurrent boundary. These are aspects of a single computational architecture, not independent functions.

5.3 Limitations

Several limitations must be acknowledged. The approximation bounds in Theorem 1 assume the bilinear form B is approximately Toeplitz on the attractor, which requires smooth spectral density. For highly non-stationary signals (e.g., language), the off-Toeplitz residual may be large. The quantitative predictions in Section 4 rely on identifying the channel activation time constant as the effective τ , which is a simplification; real channel kinetics involve multiple time constants, voltage-dependent gating, and cooperative effects. The AIS-to-auditory-neuroscience comparisons in Section 4.2–4.3 are suggestive but not definitive—the data was not collected to test this specific prediction, and confounding variables (e.g., differences in channel composition, not just AIS length) could contribute to the observed patterns.

Most critically, the claim that the AIS 'implements' Riemannian Takens attention should be understood as a mathematical correspondence, not a proven mechanism. Demonstrating that the AIS actually performs this computation requires direct measurement of the relationship between AIS structural parameters and frequency discrimination performance—ideally in a preparation where AIS length can be experimentally manipulated while frequency selectivity is measured.

5.4 Open Problems

Five problems remain open: (1) Tightening the $O(N \ln d / d)$ sidelobe bound for specific attractor geometries. (2) Formalizing the relationship between optimal τ and Lyapunov exponents. (3) Analyzing metric composition across stacked Takens layers. (4) Characterizing the gradient landscape for mosaic frequency optimization. (5) Extending the framework to non-stationary signals where the Toeplitz approximation breaks down.

6. Conclusion

We have formalized the computational primitive of the Deerskin architecture as an inner product under a Riemannian metric with Dirichlet kernel structure, proved approximation bounds for replacing standard attention with Takens-mosaic attention at dramatically reduced parameter count, and shown that the structural requirements of this mathematics map specifically to molecular features of the axon initial segment. The quantitative predictions derived from this mapping are consistent with published auditory neuroscience data.

The broader implication is that the McCulloch-Pitts neuron—and by extension all modern deep learning—operates in the flat-space limit of a fundamentally curved computation. The

weights that backpropagation laboriously learns are frozen echoes of a geometric resonance that the AIS provides through its molecular architecture. The 190 nm spectrin periodicity, the anchored channel mosaic, the activity-dependent length plasticity—these are not incidental biological details. They are the geometry that the mathematics requires.

References

- Bréchet, A. et al. (2008). Protein kinase CK2 contributes to the organization of sodium channels in axonal membranes by regulating their interactions with ankyrin G. *J. Cell Biol.* 183(6), 1101–1114.
- D'Este, E. et al. (2015). STED nanoscopy reveals the ubiquity of subcortical cytoskeleton periodicity in living neurons. *Cell Reports* 10, 1246–1251.
- Grubb, M. S. & Burrone, J. (2010). Activity-dependent relocation of the axon initial segment fine-tunes neuronal excitability. *Nature* 465, 1070–1074.
- Hedstrom, K. L. et al. (2008). AnkyrinG is required for maintenance of the axon initial segment and neuronal polarity. *J. Cell Biol.* 183(4), 635–646.
- Johnson, D. H. (1980). The relationship between spike rate and synchrony in responses of auditory-nerve fibers to single tones. *J. Acoust. Soc. Am.* 68, 1115–1122.
- Kuba, H., Ishii, T. M. & Ohmori, H. (2006). Axonal site of spike initiation enhances auditory coincidence detection. *Nature* 444, 1069–1072.
- Kuba, H., Oichi, Y. & Ohmori, H. (2010). Presynaptic activity regulates Na⁺ channel distribution at the axon initial segment. *Nature* 465, 1075–1078.
- Leterrier, C. (2018). The Axon Initial Segment: An Updated Viewpoint. *J. Neurosci.* 38(9), 2135–2145.
- Leterrier, C. et al. (2015). Nanoscale Architecture of the Axon Initial Segment Reveals an Organized and Robust Scaffold. *Cell Reports* 13, 2781–2793.
- Luode, A. & Claude (2026). The McCulloch-Pitts Neuron as a Degenerate Limit of the Deerskin Architecture. PerceptionLab preprint.
- Luode, A. & Claude (2026). The Dendrite as Translator: Signal-to-Topology Decoding in Moiré Field Networks. PerceptionLab preprint.
- Palmer, A. R. & Russell, I. J. (1986). Phase-locking in the cochlear nerve of the guinea-pig and its relation to the receptor potential of inner hair-cells. *Hearing Research* 24, 1–15.
- Takens, F. (1981). Detecting strange attractors in turbulence. In: *Dynamical Systems and Turbulence*, Lecture Notes in Mathematics 898, 366–381.
- Wang, C. et al. (2014). Bhatt structural basis for the ankyrin-mediated assembly of Nav channels. *PNAS* 111, 1–6.
- Xu, K., Zhong, G. & Zhuang, X. (2013). Actin, spectrin, and associated proteins form a periodic cytoskeletal structure in axons. *Science* 339, 452–456.